

REPRESENTATIVE PRODUCT DIAGNOSTIC

FICTIONAL PLATFORM / SAMPLE DELIVERABLE



AEGISFLOW

Making AI governance legible enough to trust.

A representative diagnostic of an AI governance platform built for regulated teams - evaluated through Clarity, Trust, Momentum, and AI Legibility.

MODEL RISK REVIEW

HUMAN APPROVAL

EVIDENCE TRACEABILITY


REGULATORY READINESS

The goal: reduce the cognitive work required to understand what the system decided, what evidence supports it, who owns the next action, and whether the workflow is actually complete.

Prepared as a portfolio sample
Not a legal compliance assessment

The platform has the right governance ingredients. The problem is that users have to assemble the meaning themselves.

AegisFlow captures policies, model risk, evidence, approvals, exceptions, and change history. But those elements are distributed across the interface in ways that make the workflow harder to read than it needs to be. The product often shows data about governance without making the governance state immediately understandable.



PRIMARY FRICTION

Too much governance is implied instead of shown.

The interface expects users to infer the relationship between AI risk classification, evidence, human sign-off, policy controls, and regulatory readiness.


RECOMMENDED DIRECTION

Turn every review into a readable decision chain.

Make the path explicit: what the system found -> why it thinks that -> what a human has confirmed -> what remains -> what happens next.

The dashboard looks complete before the decision is actually legible.

Below is a representative model review screen. The information is technically present, but three important questions are difficult to answer at a glance.



1

What does “READY” actually mean?

Controls passing is presented like an overall compliance state, while one required human approval is still pending.

2

Why is the model “HIGH” risk?

The AI judgment is prominent, but its evidence, confidence, policy mapping, and limitations are not visible where the claim is made.

3

What should I do next?

The next action is buried inside status information. The screen reads like a report, not an active governance workflow.

Diagnostic principle: A regulated workflow should not ask the user to reconcile system judgments, evidence, and human approvals mentally. The interface should carry that cognitive load.

The highest-impact gaps are not missing features. They are missing explanations between states.

1

“Regulatory readiness” overstates certainty before the review is complete.

CRITICAL

CLARITY

TRUST

AI LEGIBILITY

WHY IT MATTERS

A prominent READY state is likely to be interpreted as an overall conclusion. But the underlying workflow still contains an unresolved human approval. In a governance product, the label must distinguish control completion from review completion.

RECOMMENDED DIRECTION

Replace the global READY badge with a staged state model: Controls complete -> Human review pending -> Decision approved -> Monitoring active.

2

AI-generated risk classification is visible; its basis is not.

CRITICAL

TRUST

AI LEGIBILITY

WHY IT MATTERS

“HIGH - AI generated” creates an authoritative judgment without the context needed to evaluate it. Users need to distinguish model-derived assessment from policy requirements and human decisions.

RECOMMENDED DIRECTION

Expose a compact decision basis beside the classification: triggering factors, source evidence, policy/control mappings, uncertainty, and an explicit “system recommendation - human confirmation required” state.

3

System checks and human approvals are visually mixed into one progress state.

HIGH

CLARITY

MOMENTUM

AI LEGIBILITY

WHY IT MATTERS

A user can see “2 of 3,” but not whether the missing step is an automated test, a policy owner, legal review, or executive approval. That ambiguity slows handoffs and makes ownership unclear.

RECOMMENDED DIRECTION

Separate machine state from human state. Show automated assessment complete, named reviewer pending, due date, blocker, and one explicit next action.

The secondary gaps compound the same problem: evidence exists, but action and rationale are hard to read.

4

Failed controls explain the problem but do not create a clear remediation path.

HIGH

MOMENTUM

CLARITY

WHY IT MATTERS

A failed control currently sends users into evidence and policy tabs to determine what will clear the failure. That forces users to translate a finding into a task themselves.

RECOMMENDED DIRECTION

Turn each unresolved control into an actionable remediation object: owner, required evidence, acceptance condition, due date, and re-check action.

5

The audit trail records events, not the rationale behind the decision.

HIGH

TRUST

AI LEGIBILITY

WHY IT MATTERS

"Risk changed from Medium to High" is insufficient when someone later needs to understand why. The durable artifact is not just the event - it is the evidence and reasoning snapshot that made the decision defensible at that moment.

RECOMMENDED DIRECTION

Capture decision snapshots: prior state, new state, triggering evidence, AI recommendation, human override/confirmation, rationale, and policy version.

6

The dashboard prioritizes inventory over unresolved governance work.

MEDIUM

MOMENTUM

CLARITY

WHY IT MATTERS

Counts such as "18 evidence files" describe the system but do not help a governance lead decide what to do today. The homepage should surface incomplete or aging decisions before passive inventory.

RECOMMENDED DIRECTION


Lead with decision queues: reviews awaiting approval, evidence gaps, expired controls, unresolved exceptions, and

PATTERN ACROSS ALL SIX FINDINGS


AegisFlow does not need more governance information. It needs stronger relationships between information, judgment, ownership, and action.

Make every model review readable as one decision chain.

The future-state experience should make it possible to answer five questions without navigating away from the review.



FUTURE-STATE WORKFLOW



AI legibility standard: whenever AegisFlow makes or influences a judgment, the user should be able to see what the system concluded, the evidence it used, the limits of that conclusion, and where human authority begins.

Fix the decision layer before expanding the feature layer.

The order matters. The first changes should make existing governance logic readable; only then should AegisFlow add more automation, scoring, or reporting.

PRIORITY 01

Reframe the review around a single decision state

HIGH IMPACT

Replace ambiguous "readiness" with explicit control, human review, approval, and monitoring states. Make incomplete work impossible to mistake for a finished decision.

CLARITY + TRUST

PRIORITY 02

Add an AI decision-basis component

HIGH IMPACT

For AI-generated classifications, expose the factors, evidence, policy relationship, uncertainty, and human confirmation requirement at the point of judgment.

AI LEGIBILITY + TRUST

PRIORITY 03

Separate machine completion from human ownership

HIGH IMPACT

Show which system checks are complete, who owns the next human step, what is blocking them, and the single next action.

MOMENTUM + CLARITY

PRIORITY 04

Turn findings into remediation objects

MEDIUM-HIGH

Every failed control or missing evidence item should state what resolves it, who owns it, and how the platform will recognize completion.

MOMENTUM

PRIORITY 05

Upgrade audit history into decision history

MEDIUM-HIGH

Preserve evidence and rationale snapshots so a reviewer can reconstruct why a decision was made at the time.

TRUST + AI LEGIBILITY

WHAT THIS CHANGES FOR THE USER

Less reconstruction. The user no longer has to mentally connect policies, evidence, model judgments, reviewer status, and next actions.

WHAT THIS CHANGES FOR THE PRODUCT

Governance becomes legible. The platform stops looking like a database of controls and starts behaving like a decision system people can understand and defend.

The opportunity is not to make AegisFlow feel simpler than governance is. It is to make the complexity readable.

Regulated AI workflows will always contain policy, evidence, uncertainty, human judgment, exceptions, and monitoring. The interface earns trust when those relationships are visible enough that a user can understand the decision without reconstructing the system in their head.

CLARITY

Show what state the review is actually in.

TRUST

Make evidence and decision rationale visible at the point of judgment.

MOMENTUM

Turn unresolved governance work into clear ownership and next actions.

AI LEGIBILITY

Make the system's role, reasoning inputs, uncertainty, and human boundary explicit.

REPRESENTATIVE ENGAGEMENT PATH

DIAGNOSTIC

PRIORITIZED DIRECTION

IMPLEMENTATION SCOPE

BUILD

If the problem is already known and the direction is clear, the diagnostic can be skipped and implementation can begin from a defined scope.

About this sample: AegisFlow is fictional. This document demonstrates the structure and level of thinking of a representative product diagnostic. It evaluates product communication, workflow legibility, and interaction design - not legal compliance, regulatory sufficiency, or legal advice.